

```
!pip install -U transformers
```

Show hidden output

```
!pip install -U datasets
```

Show hidden output

Local Inference on GPU

Model page: <https://huggingface.co/casperhansen/llama-3-8b-instruct-awq>

⚠️ If the generated code snippets do not work, please open an issue on either the [model repo](#) and/or on [huggingface.js](#) 🙋

```
!pip install gptqmodel
```

Show hidden output

```
!uv pip install vllm==0.18.0 --torch-backend=auto
```

Show hidden output

```
# Library imports
from vllm import LLM, SamplingParams
```

WARNING:torchao:Skipping import of cpp extensions due to incompatible torch version. Please upgrade to torch >= 2.11.0 (found 2.

```
# Block to fetch prompts from the alpaca dataset
from datasets import load_dataset
ds = load_dataset("tatsu-lab/alpaca", split = "train")
inputs = ds["instruction"][:200]
```

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF\_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>), set it
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.

warnings.warn(

README.md: 7.47k/? [00:00<00:00, 476kB/s]

data/train-00000-of-00001-a09b74b3ef9c3b(...): 100%

24.2M/24.2M [00:01<00:00, 21.5MB/s]

Generating train split: 100%

52002/52002 [00:00<00:00, 219336.75 examples/s]

```
# Initialize the vLLM object
llm = LLM(model="casperhansen/llama-3-8b-instruct-awq", quantization = "awq", tensor_parallel_size = 1)
```

Show hidden output

```
# Set parameters for the output
params = SamplingParams(temperature=0.8, max_tokens = 100)
```

```
# Get the tokenizer for applying chat template
tokenizer = llm.get_tokenizer()
```

```
formatted_prompts = []
for prompt in inputs:
    # Wrap the raw text in the standard dictionary format
    message = [{"role": "user", "content": prompt}]

    # Apply the template to create the raw string Llama-3 expects
    formatted_string = tokenizer.apply_chat_template(
        conversation=message,
        tokenize=False,
        add_generation_prompt=True
    )

    formatted_prompts.append(formatted_string)
```

```
outputs = llm.generate(formatted_prompts, sampling_params = params)

for output in outputs:
    prompt = output.prompt
    generated_text = output.outputs[0].text
    print(f"Prompt: {prompt!r}, Generated text: {generated_text!r}")
```

[Show hidden output](#)

Start coding or [generate](#) with AI.